

MAVIS: VISUALIZATION TOOL FOR MULTIPLE SEQUENCE ALIGNMENTS

A Thesis

Presented to

the Faculty of the Department of Computer Science

University of Houston

In Partial Fulfillment

of the Requirements for the Degree

Master of Science

By

Lei Zhao

May 2011

MAVIS: VISUALIZATION TOOL FOR MULTIPLE SEQUENCE ALIGNMENTS

Lei Zhao

APPROVED:

Dr. Yuriy Fofanov, Chairman
Department of Computer Science

Dr. Dan Graur
Department of Biology and Biochemistry

Dr. Shishir Shah
Department of Computer Science

Dean, College of Natural Sciences and Mathematics

MAVIS: VISUALIZATION TOOL FOR MULTIPLE SEQUENCE ALIGNMENTS

An Abstract of a Thesis
Presented to
the Faculty of the Department of Computer Science
University of Houston

In Partial Fulfillment
of the Requirements for the Degree
Master of Science

By
Lei Zhao
May 2011

Abstract

Write your abstract

Contents

1	Introduction	1
1.1	Background	1
1.2	Related Work	2
1.3	Thesis Outline	3
2	Approach	4
2.1	Data Input	4
2.2	Color Generating	5
2.3	Color Optimization	6
2.4	MSA Visualization	6
3	Algorithms	7
3.1	Multidimensional Scaling	7
3.2	Color Space Mapping	9
3.3	Hue Rotation and Optimization	10
3.4	Hue Flipping	11
3.5	Sequence Sorting	12
4	Implementation	13
4.1	R Script	14
4.2	API	15

4.3	User Interface	15
5	Results	16
5.1	Validity Test	16
5.2	Performance Test	17
6	Discussion	18
	References	21

List of Figures

List of Tables

Chapter 1

Introduction

1.1 Background

Multiple sequence alignment (MSA) analysis, a fundamental approach in molecular biology, is required for nearly all comparative sequence studies, including evolutionary, functional and structural aspects. A multiple sequence alignment is usually a matrix, in which each row represents a sequence, and each column corresponds to the equivalent positions across all sequences aligned. Each element of the matrix contains a symbol, which can be either '-' (a gap) or a sequence symbol (an amino acid for DNA/RNA sequences, or a nucleotide base for protein sequences). [1]

Due to the size and complexity of the multiple sequence alignment datasets, computer visualization techniques are essential to assist in understanding, analyzing and evaluating the alignment results. Numerous MSA visualization applications,

either stand-alone or web-based, have been developed in recent decades, ranging from simple MSA viewers, to complex editing and analysis platforms. Many of them have graphical user interfaces showing a grid of symbols, usually along with various coloring, sizing and annotation. [2]

1.2 Related Work

Coloring an alignment has been widely implemented to highlight specific regions, properties or patterns. [2] The simplest way of coloring is to use a fixed color scheme, in which each sequence symbol has its own color, and will not be changed across the alignment. The assignment of color schemes are usually based on some empirical properties or chemical classifications. There are numerous color schemes available for different purposes, but Taylor [3] and Clustal [4] are widely used and often considered de facto standards.

However, using pre-determined color schemes to color sequence symbols, either amino acids or nucleotide bases, are not always helpful to phylogenetic analysis. Those fixed schemes focus on chemical differences between specific types of symbols, rather than between sequences, regions, blocks or structures of the alignment. For example, if symbols are examined within a certain column, the colors reflect the similarities fairly well. However, if the alignment is analyzed horizontally, comparative information within a row or a rectangular region are unable to be displayed, like how one sequence is similar with others, which section aligns better than another, which subset of sequences in which range of columns are more conserved, or how robust is

the alignment accuracy.

Here we suggest a new MSA visualization tool, Mavis (Multiple Alignment Visualization), which focuses on highlighting evolutionary relationships and internal structures of an alignment. Instead of relying on chemical properties or any other information from symbol types, our approach takes only symbol-symbol correlation into consideration. The most valuable result we would like to provide is a whole picture of the MSA.

1.3 Thesis Outline

This thesis is organized in the following way: Chapter 2 briefly summarizes how Mavis processes data input, creates colors and displays the colored MSA graph. Chapter 3 explains the algorithm step by step, from scaling to mapping and optimization. Chapter 4 describes our implementation of the algorithm, including the back-end, API and front-end. Chapter 5 presents test results to show the validity and performance of our implementation. The whole work is discussed and concluded in Chapter 6.

Chapter 2

Approach

In this thesis, we address a different method to visualize multiple sequence alignments. The main objective is to correlate alignment similarity (or confidence) with color similarity. In the alignment matrix, a well-aligned area forms a solid-color shape, while a poor-aligned one becomes a mosaic of various colors. Alignment-level information, such as evolutionary relationships, phylogenetic characteristics and alignment accuracy should be revealed by examining the graph generated by our tool.

2.1 Data Input

Unlike most traditional coloring approaches, we start with an set of pairwise scores, rather than the alignment itself. Each of these scores represents the similarity or alignment confidence between two symbols aligned together in the same position, or,

in the same column. Every pair of symbols in the same column has a score.

This set of data can be generated by GUIDANCE [5, 6], a measure quantifying alignment uncertainty. It produces a confidence score between 0 and 1, named GUIDANCE score, for each symbol, symbol pair, column and sequence in the alignment. The symbol-pair GUIDANCE confidence scores can be used in our study to visualize MSAs.

2.2 Color Generating

The colors of the symbols are generated in such way that the color differences represent symbol differences. We consider the confidence scores reversely, as distances, and convert all the scores of a column into a $N * N$ distance matrix, where N is the number of symbols. These N symbols can be represented by N points in at most $N - 1$ dimensions, such that the Euclidean distance between any pair of points is exactly equal to the distance in the matrix. Once we can map these points to a color space, those corresponding colors should be able to fulfill our requirements.

However, most color spaces are three-dimensional or under. It is necessary to reduce the dimension from $N - 1$ to at most 3 in order to process the mapping, but still preserve the distances as much as possible. This dimension reduction is done by a statistical technique called *classical multidimensional scaling* [7].

2.3 Color Optimization

Up to now, our approach only operates on the column level. Further adjustment is still needed to achieve the alignment-level objectives. We notice that the color space can be arbitrarily rotated and/or flipped, without changing the distance between any pair of symbols. For instance, all reds and all greens are virtually identical in the sense of distances and quality of alignment. Converting them into the same color pattern will better support this idea and remove unnecessary color noises.

We defined a penalty function to quantify the overall noise level of an alignment. The ideal combination of rotating and flipping each column should result in the lowest penalty value. For large MSAs, finding the global optimum is not always possible, and the approximate solution is also an acceptable option.

2.4 MSA Visualization

We choose to implement a web-based user interface, because of its flexibility, portability and light weight. Many features are provided to enhance the visualization effects, such as sorting sequences by their average colors, toggling to show/hide symbol names, and multiple-MSA support. An API between the front-end and the back-end is designed to make the system more extensible.

Chapter 3

Algorithms

The algorithm of color computation consists of five stages. First, multidimensional scaling reduces the number of dimensions and convert the confidence/similarity scores into coordinates. Second, coordinates are mapped onto part of the CIE Lab color space[8], and colors are created. Third, colors are flipped and rotated to smooth the graph and remove color noises. Finally, calculate the average color/hue of each sequence, sort them, and output for further display.

3.1 Multidimensional Scaling

This stage begins with the confidence scores described above. Each column in the MSA has a complete set C of pairwise scores. Each pair of non-gap symbols i, j in this column has a score $m_{ij} \in [0, 1]$.

$$M = 1$$

This similarity score S can be easily converted to a dissimilarity score, or a distance D , between the two symbols, by $D = 1 - S$. All these distances within a column can be re-organized as a symmetric distance matrix.

We try to map all the symbols to a two-dimensional space, in which they still preserve their distances as much as possible. This procedure involves a statistical technique called classic multidimensional scaling, also known as Torgerson-Gower Scaling or Principle Coordinates Analysis (Gower, 1966). Multidimensional scaling takes a distance matrix and assigns a coordinate to each item in a lower dimensional space, such that the Euclidean distance between any two coordinates is approximately equal to the dissimilarity score between the corresponding items.

An R function called `cmdscale()` is utilized in the multidimensional scaling procedure. The function is written in R, which is a programming language for statistical computing (Ihaka and Gentleman 1996), and is provided in Rs stat package. It reads a distance matrix and returns a set of points in a k -dimensional space, where k is a user-defined parameter to the function, and must be less than the number of points. (Becker et al. 1988, Cailliez 1983, Cox and Cox 2001)

Here is our approach, k is set to 2, meaning the distance matrix should be scaled to a two-dimensional space. However, a column in a multiple sequence alignment, which represents an equivalent position across all the sequences, can contain only one or two non-gap symbols, so that the column has only one or two points. In such cases, $k=2$ is not a valid parameter because k must be less than the number of points, so k should be set to zero (for one point) or one (for two points), and the missing coordinates should be filled by zeros.

The input data set represents the similarities between symbols, while the multi-dimensional scaling procedure requires the dissimilarities (or distances). We simply convert similarity (S) to dissimilarity (D) by $D = 1 - S$ and pass the converted matrix to the `cmdscale()` function.

3.2 Color Space Mapping

The multidimensional scaling procedure returns a set C of coordinates on a two-dimensional space, representing the symbols in a multiple sequence alignment. We further map this two-dimensional space into an three-dimensional CIE Lab color space.

The CIE Lab color space is a color-opponent space. The name Lab stands for the three dimensions: L for lightness, and a and b for the color opponents based on nonlinearly compressed CIE XYZ color space coordinates. Compared to the RGB and CMYK color models, the design of Lab color emphasizes more on approximating human vision rather than physical devices. This feature makes it suitable for our visualization purpose.

Since the data set C from the previous procedure contains only two-dimensional coordinates, and we only use the most visible colors instead of the whole color space, the dimension L for lightness is set to 75. The two dimensions from data set C are scaled from $[0,1]$ to $[0, 100]$ and used as color opponents a and b.

3.3 Hue Rotation and Optimization

One of the limitations in the previous multiple sequence alignment visualization techniques is the unnecessary color difference among columns, which doesn't represent any biological dissimilarity but introduces much confusion to the analysis. To minimize this effect, we convert the colors from CIE Lab space to CIE LCH space, perform rotations and minimize the the sum of hue variances of all the sequences. This procedure eliminates between-column color differences to the greatest extent possible, smooth the color pattern and emphasize the real similarities and dissimilarities among sequences and blocks.

The CIE LCH color space, also known as polar-Lab color space, is a transformation of the CIE Lab space so that the a and b axes are converted to a polar coordinate system, which makes it easier to rotate. The radial coordinate C measures chroma and the angular coordinate H measures hue. The CIE LCH colors can be converted from CIE Lab colors using an R library called `colorspace`.

The fact that 0 and 2 are identical angles, so that given two angles p and q, the difference between them can be calculated by $\text{diff}(t1, t2) = (p - q + \pi) \bmod 2\pi$. This gives the difference in the interval $[-\pi, \pi)$, or in degrees $[-180, 180)$. The absolute value of the difference is in the interval $[0, \pi]$, or in degrees $[0, 180]$. Given a set A of angles $a[i]$, $i \in [0, n-1]$, the average angle $\text{avg}(A)$ is

and the variance $\text{var}(A)$ is

Given an multiple sequence alignment M, with m rows and n columns, we look for a optimized set R of n rotation angles $r[i]$, such that after rotating each hue in the

ith column of the alignment clockwise by $r[i]$, then the penalty function is minimized.

The optimization procedure is performed with R's general-purpose optimization function `optim()`. Among five available optimization algorithms, we choose the bounded Broyden-Fletcher-Goldfarb-Shanno (L-BFGS-B) method. This algorithm allows box constraints, meaning each variable can be given a lower and/or upper bound. Plus, it uses a limited amount of memory. The initial values are set to 0 and boundaries to $[0, 2]$. More detailed comparison between different optimization functions and algorithms will be discussed later in Chapter 6.

3.4 Hue Flipping

One characteristic of the coordinates generated from multidimensional scaling is, they can not only be rotated, but also flipped, without changing the distances from each other. Adding flipping ability to the above rotation optimization procedure may result in lower penalty values and better coloring results. To keep the penalty function simple and fast, we perform an heuristic flipping before the rotation step.

For each i th column ($2 \leq i \leq n$), compare it to the previous one. Pick rows which exists in both columns and record the hues in two arrays H1 and H2. For each of the two arrays, compare each hue to the previous one using the `diff()` function described above. For each `diff()` $\geq \pi/8$, record a clockwise change; for each `diff()` $\leq -\pi/8$, record a counterclockwise change; otherwise leave it as an insignificant change. Within a column, if the number of clockwise changes is greater than the number of counterclockwise ones, call it a clockwise column, and vice versa. If two adjacent

columns are in opposite types, flip the latter one to make them in same type.

After comparisons of all adjacent column pairs have been finished, no two pairs will be in opposite types, and the flipped data set is passed to the optimization procedure described above.

3.5 Sequence Sorting

The alignment graph is supposed to provide information on how sequences will group with each other. We sort sequences after coloring optimization, based on the average hue of each one, which is calculated by the average angle function `avg()` proposed previously.

Since the hue value is circular, where to break the sorted circle is a question. In the circle, we look for the biggest gap without any points in it, and break it at the midpoint. This keeps similar sequences together as much as possible.

Chapter 4

Implementation

We implemented the whole visualization pipeline with an R script as back-end, a Perl CGI script as API, and an HTML/JavaScript web page as front-end.

The R script `coloring.r` takes the pair-wise similarity score file as input, performs multidimensional scaling and color optimization, sorts sequences, and output colors and sequence information into two tab-separated files `colors.tsv` and `seqinfo.tsv`. The front-end web page takes an alignment ID as parameter and sends it to the API. The Perl CGI script parse the color file and sequence information file, along with the FASTA-format alignment sequences data, and sends them back to front-end in JSON format. The front-end uses JavaScript to generate the alignment matrix and render it with colors and symbol letters.

4.1 R Script

The script `coloring.r` begins with reading three arguments from command-line calls: score file name, colors file name and sequence information file name. The first file is for input and the latter two are for output. An example of calling this script on UNIX operating system is given below. `Rscript coloring.r data/score.tsv data/colors.tsv data/seqinfo.tsv`

The file `score.tsv` contains four columns separated by tabs: column number, row number one, row number two, and similarity score between these two rows in this column. The row and column numbers start from one, and the similarity score values range from zero to one.

The script `coloring.r` consists of four parts. The first part reads the input file, performs multidimensional scaling and create color for each symbol in the CIE LCH space. The second part rotates and flips color hues and optimize the penalty function. The third part sorts the sequences by their average hues and calculates each sequences average RGB color. The last part outputs the symbol colors and sequence information into corresponding files.

The color file `colors.tsv` has five columns separated by tabs: column number, row number, and the symbol color in RGB triplet (red, green, blue) in the range from 0 to 1, which will be converted into hexadecimal format used in HTML and JavaScript. The sequence information file `seqinfo.tsv` also has five columns: row number, average hue from 0 to 360 degrees, average color in RGB triplet.

4.2 API

A Perl CGI script `api.pl` and a Perl package `Mavis::API` act as the API between the R back-end and HTML/JavaScript front-end. The Perl CGI `api.pl` takes two parameters from HTTP requests, `action` and `id`. The default action is `alignment`, meaning the alignment information, including sequences, colors, average colors and order of sequences. In this action, an additional `id` parameter must be provided as the alignment ID. Another available action is `list`, which requests all existing alignment IDs.

The API returns the query result to the web page in JSON (JavaScript Object Notation) format, which is a lightweight text-based data-interchange standard. JSON is derived from a subset of JavaScript syntax and can be easily parsed in JavaScript as well as many other languages, like Perl. We included the Perl JSON 2.51 module to provide JSON encoding/decoding ability to our API.

4.3 User Interface

We implement an web interface with HTML and JavaScript to send API requests and draw visualization graphs. The JavaScript code dynamically communicates with API, creates an HTML table, and fills in colors and symbols. This part is simplified by jQuery, which is a popular cross-platform JavaScript library that abstracts complicated DOM selection, CSS manipulation and Ajax effects.

Chapter 5

Results

To evaluate the validity and performance of our approach, we tested it in two different ways. First, we run it on a few artificial datasets, to see if it creates the color patterns as we expected. Then, we run it on the BALiBASE dataset, observe how alignment size, in both number of rows and columns, will affect the execution time, and compare it with other tools.

5.1 Validity Test

(to do)

5.2 Performance Test

(to do)

Chapter 6

Discussion

In this report, we introduced a new way of coloring and visualizing multiple sequence alignments based on symbol-symbol similarity scores. Our main objective has been to create an intuitive way of showing the quality and internal structure of an alignment, using colors, which is the most natural and sensible way for human eyes. The basic idea is, the greater the dissimilarity between sequences, the more obvious difference in colors.

Most previous techniques were based on fixed color schemes, meaning each sequence symbol, like an amino acid or a nucleotide, is assigned to a pre-selected color. This is of course a simple and straightforward solution, and is easy for observers to find a particular symbol or pattern. However, this fixed-scheme approach does not emphasize the relationship between adjacent columns and the internal regions and structures in the level of the whole alignment. This is the main reason why we don't use any predetermined color scheme, but calculate colors only based on the

distance matrices, and rotate and flip colors to make them as smooth as possible. This strategy brings significant improvement to the alignment coloring, by showing the conserved regions and blocks in same or similar colors, while those irrelevant ones in quite different colors.

Since we are to convert the scaled coordinates to colors, choosing a suitable color space is an fundamental task. At first we scaled the distance matrix down to a three-dimensional space and mapped it directly to the RGB color space. However, we soon found a problem, that some colors, like very dark ones and very light ones, did not perfectly serve the purpose of representing distances, yet made the graph more noisy. We realized that we dont really need the whole color space and all possible colors. Instead, those colors with proper range of lightness and different hues will be enough to do the job. So we chose to scale down to a two-dimensional space and map it to CIE Lab space with a fixed lightness value (75). The reason why we didnt choose one-dimensional scaling is that, a linear space either could not map to enough number of colors (for example, use only one primary color, like red), or could not preserve the distance information (for example, use only hues). So two dimensions is a good balance.

In R, there are several general purpose optimization packages which offer facilities for solving our color rotation and flipping problems. Two popular functions are `optim()` and `nlminb()` from package `stats`. Function `optim()` provides implementations of five algorithms: Broyden-Fletcher-Goldfarb-Shanno (BFGS), bounded BFGS (L-BFGS-B), conjugate gradient (CG), Nelder and Mead (Nelder-Mead), and simulated annealing (SANN). Nelder-Mead (Nelder and Mead 1965), which is the default one,

returns robust results but is relatively slow. CG (Fletcher and Reeves 1964) is faster in larger optimization problems, but more fragile. BFGS is a balance, and L-BFGS-B further provides the ability of box constraints, that each variable can be given a lower/upper bound. SANN (Belisle 1992) is more powerful on rough surfaces but relatively slow. Another function `nlminb()` offers similar box constraint optimization and similar performance to L-BFGS-B, so these two algorithms are chosen to a further test.

Optimization algorithms always suffer from the local versus global minimum problem, and the final result are more or less unstable and depending on the initial values. To decide which one of L-BFGS-B and `nlminb()` is more stable in our approach, we run a test on both of them. The test dataset is the alignment of 44 Vpu protein sequences from Guidance. We randomly created 100 sets of initial values, performed optimizations, and see how the return value of the penalty function described in section 3.3 changed.

Algorithm	Minimum penalty	Average penalty	Maximum penalty	Standard deviation
Initial	370,584	425,728	444,771	13,998
L-BFGS-B	136,909	178,850	247,218	17,871
L-BFGS-B - run 3 times	136,909	178,850	247,218	17,871
<code>nlminb()</code>	146,759	181,688	426,535	34,579
<code>nlminb()</code> - run 3 times	142,888	171,380	325,469	21,402

Obviously using L-BFGS-B is more stable than `nlminb()` on our test dataset.

References

- [1] Robert C Edgar and Serafim Batzoglou. Multiple sequence alignment. *Current Opinion in Structural Biology*, 16(3):368–373, June 2006.
- [2] JB Procter, J Thompson, I Letunic, C Creevey, F Jossinet, and GJ Barton. Visualization of multiple alignments, phylogenies and gene family evolution. *Nat Methods*, 7(3 Suppl):S16–25, 2010.
- [3] Kuang Lin, Alex C. W. May, and William R. Taylor. Amino acid encoding schemes from protein structure alignments: Multi-dimensional vectors to describe residue types. *Journal of Theoretical Biology*, 216(3):361 – 365, 2002.
- [4] Julie D Thompson, Toby. J Gibson, and Des G Higgins. Multiple sequence alignment using ClustalW and ClustalX.
- [5] Osnat Penn, Eyal Privman, Giddy Landan, Dan Graur, and Tal Pupko. An alignment confidence score capturing robustness to guide tree uncertainty. *Molecular Biology and Evolution*, 27(8):1759 –1767, 2010.
- [6] O. Penn, E. Privman, H. Ashkenazy, G. Landan, D. Graur, and T. Pupko. GUID-ANCE: a web server for assessing alignment confidence scores. *Nucleic Acids Research*, 38(Web Server):W23–W28, May 2010.
- [7] Ingwer Borg and Patrick J. F. Groenen. *Modern multidimensional scaling: theory and applications*. Springer, 1997.
- [8] K. McLaren. XIII—The development of the CIE 1976 (L^* a^* b^*) uniform colour space and colour difference formula. *Journal of the Society of Dyers and Colourists*, 92(9):338–341, September 1976.